

Linux kernel and C programming

BLOCK 2: Linux kernel modules

Gergely Korcsák

Óbuda University
korcsak.gergely@nik.uni-obuda.hu

March 11, 2025

Presentation Overview

- 1 C Leftover shorts
 - goto
- 2 Linux kernel modules
 - Introduction
 - Headers
 - Hello World
 - Building
 - Inclusion
 - Init module with macros
 - Further module description
 - Proc files
- 3 Exercises
- 4 References

C Leftover shorts

goto

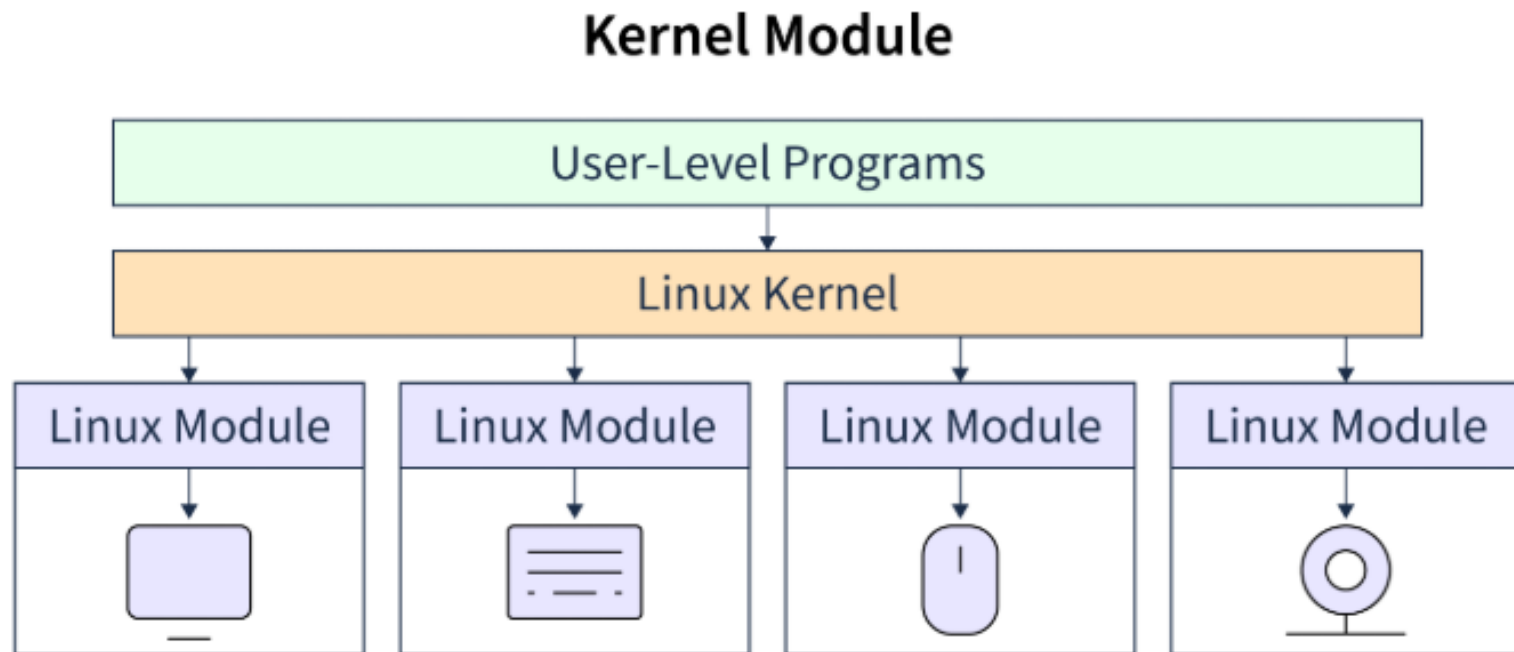
goto

With the `goto` keyword, we can jump to a specific line to continue execution from.

```
1 int main() {  
2     printf("hello 1");  
3  
4     goto section2;  
5  
6     printf("hello 2"); // Will not execute  
7  
8 section2:  
9     printf("hello 3");  
10  
11     return 0;  
12 }
```

Linux kernel modules

Introduction



Headers

To build modules, first we will need to get the kernel headers.

```
1 $ sudo apt install build-essential kmod \  
2     linux-headers-$(uname -r)
```

With which we can now access the kernel functions, with dynamic linking.

Hello World

We can create our first module:

```
1 // hello_module.c
2 #include <linux/module.h> /* Needed by all modules */
3 #include <linux/printk.h> /* Needed for pr_info() */
4
5 int init_module(void)
6 {
7     pr_info("Hello world\n");
8
9     return 0;
10 }
11
12 void cleanup_module(void)
13 {
14     pr_info("Goodbye world\n");
15 }
16
17 MODULE_LICENSE("GPL");
```


Building

We can build it with the following `Makefile`

```
1 # Makefile
2 # We need to add the C files as object files
3 obj-m += hello_module.o
4
5 BUILDDIR := $(shell pwd)/build
6 SRCDIR   := $(shell pwd)
7 KERNELDIR := /lib/modules/$(shell uname -r)/build
8
9 all:
10     $(MAKE) -C $(KERNELDIR) M=$(BUILDDIR) src=$(SRCDIR)
11     modules
12
13 clean:
14     $(MAKE) -C $(KERNELDIR) M=$(BUILDDIR) src=$(SRCDIR)
15     clean
```

Then build with `make all`.

Inclusion

modinfo

We can check the module information:

```
1 $ modinfo build/hello_module.ko
2 filename:          /.../build/hello_module.ko
3 version:           0.1
4 description:
5 author:            Bob
6 ...
```

Inclusion

lsmod

Show the currently running modules

```
1 $ lsmod
2 ahci                23756      3
3 libahci             65734      1 ahci
4 sha1_ssse3          35664      1
5 ...
```

Inclusion

insmod

Include a new module by path

```
1 $ sudo insmod build/hello_module.ko
```

Lets check it with lsmod

```
1 $ lsmod | grep hello
2 module_hello      12288  0
```

rmmod

And remove it with

```
1 $ sudo rmmod hello_module
```

dmesg

We can check the output of a module in `dmesg`

```
1 $ dmesg | tail -2
2 [ 2645.234345] Hello World!
3 [ 2646.234345] Goodbye World!
```

Init module with macros

We need to modify our code a bit. Lets include the `init.h` header as well, and specify the init functions by macros:

```
1 // ...
2 #include <linux/init.h>
3 // ...
4
5 // Modify init_module(void) to
6 static int __inti hello_init(void) {
7     // ...
8
9 // And cleanup_module(void) to
10 static void __exit hello_exit(void) {
11     // ...
12
13 // And add module_init(...) and module_exit(...)
14 module_init(hello_init);
15 module_exit(hello_exit);
```

Further module description

To build a module, we must define its license with `MODULE_LICENSE("GPL");` for example. But we can also specify the author and the description with:

```
1 MODULE_LICENSE("GPL");  
2 MODULE_AUTHOR("LKMPG");  
3 MODULE_DESCRIPTION("A sample driver");
```

Lets add our description, and check it with `modinfo`.

Proc files

Proc files were originally created to access information. They are located at the `/proc` folder. These files can be created when loading a module, and its the modules job to remove them afterwards. Proc also has some useful systemfiles, like `meminfo` and `cpuinfo`.

Proc files

To create our own proc file, we need to include:

```
1 #include <linux/proc_fs.h>
2 #include <linux/uaccess.h>
```

Then fill create our functions for:

```
1 // coming from proc_fs.h
2 struct proc_ops {
3     ssize_t (*proc_read)(struct file *, char __user *,
4                          size_t, loff_t *);
5     ssize_t (*proc_write)(struct file *,
6                          const char __user *,
7                          size_t, loff_t *);
8     int (*proc_open)(struct inode *, struct file *);
9     int (*proc_release)(struct inode *, struct file *);
10    // ...
11 };
```

Proc files

Then crate our own struct

```
1 // create our own struct
2 static struct proc_ops my_proc_file {
3     .proc_read = my_proc_read,
4     .proc_write = my_proc_write,
5 };
```

Proc files - read

```
1 static ssize_t my_proc_read(  
2     struct file* fp, char __user* buffer,  
3     size_t buffer_length, loff_t* offset)  
4 {  
5     char s[] = "Hello World!\n";  
6     if (*offset >= sizeof(s) ||  
7         copy_to_user(buffer, s, sizeof(s)))  
8     {  
9         pr_info("copy_to_user failed\n");  
10        return 0; // copy failed, written 0 bytes  
11    }  
12    else  
13    {  
14        pr_info("procfile read %s\n",  
15            fp->f_path.dentry->d_name.name);  
16        *offset += sizeof(s);  
17        return sizeof(s);  
18    }  
19 }
```

Proc files - write

```
1 static ssize_t my_proc_write(  
2     struct file* fp, char __user* buffer,  
3     size_t length, loff_t* offset)  
4 {  
5     char copy_buf[100];  
6     size_t cp_size = length;  
7     if (cp_size > sizeof(copy_buf))  
8         cp_size = sizeof(copy_buf) - 1;  
9  
10    if (copy_from_user(copy_buf, buffer, cp_size))  
11        return -EFAULT;  
12  
13    copy_buf[cp_size] = '\n';  
14    *offset += cp_size;  
15    pr_info("procfile write: %s\n", copy_buf);  
16  
17    return cp_size;  
18 }
```

Proc files

Then we can create the file with

```
1 static struct proc_dir_entry* our_proc_file;
2
3 // in init(..)
4 our_proc_file = proc_create("hello", 0644, NULL,
5                             &my_proc_file);
6 if (NULL == our_proc_file) {
7     pr_alert("Error: Could not initialize /proc/hello\n");
8     return -ENOMEM;
9 }
```

And remove it with

```
1 // in exit(..)
2 proc_remove(our_proc_file);
```

Excercises

- 1 Lets create our first hello_module example
- 2 Create a proc file, to which we can write to, and read out its value. If we are out of buffer space, then return an error.

References

- 1 <https://sysprog21.github.io/lkmpg/>
- 2 <https://elixir.bootlin.com/linux/v6.13.5/source/include/linux/>